

Homework 2

PSTAT Summer 2024

Due date: July 15th, 2024 at 23:59 PT

1. This question uses the *cereal* data set available in the Homework Assignment 2 on Canvas. The following command can be used to read the data into R. Make sure the “cereal.txt” file is in the same folder as your R/Rmd file.

```
Cereal <- read.table("cereal.txt",header=T)
str(Cereal)
```

```
## 'data.frame':    77 obs. of  16 variables:
## $ name      : chr  "100%_Bran" "100%_Natural_Bran" "All-Bran" "All-Bran_with_Extra_Fiber" ...
## $ mfr       : chr  "N" "Q" "K" "K" ...
## $ type      : chr  "C" "C" "C" "C" ...
## $ calories: int  70 120 70 50 110 110 110 130 90 90 ...
## $ protein  : int  4 3 4 4 2 2 2 3 2 3 ...
## $ fat       : int  1 5 1 0 2 2 0 2 1 0 ...
## $ sodium   : int  130 15 260 140 200 180 125 210 200 210 ...
## $ fiber    : num  10 2 9 14 1 1.5 1 2 4 5 ...
## $ carbo    : num  5 8 7 8 14 10.5 11 18 15 13 ...
## $ sugars   : int  6 8 5 0 8 10 14 8 6 5 ...
## $ potass   : int  280 135 320 330 -1 70 30 100 125 190 ...
## $ vitamins: int  25 0 25 25 25 25 25 25 25 ...
## $ shelf    : int  3 3 3 3 3 1 2 3 1 3 ...
## $ weight   : num  1 1 1 1 1 1 1 1.33 1 1 ...
## $ cups     : num  0.33 1 0.33 0.5 0.75 0.75 1 0.75 0.67 0.67 ...
## $ rating   : num  68.4 34 59.4 93.7 34.4 ...
```

The data set *cereal* contains measurements for a set of 77 cereal brands. For this assignment only consider the following variables:

- Rating: Quality rating
- Protein: Amount of protein.
- Fat: Amount of fat.
- Fiber: Amount of fiber.
- Carbo: Amount of carbohydrates.
- Sugars: Amount of sugar.
- Potass: Amount of potassium.
- Vitamins: Amount of vitamins.
- Cups: Portion size in cups.

Our goal is to study how *rating* is related to all other 8 variables.

- (a) (4pts) Explore the data and perform a descriptive analysis of each variable, include any plot/statistics that you find relevant (histograms, scatter diagrams, correlation coefficients). Did you find any outlier? If

yes, is it reasonable to remove this observation? why?

(b) (3pts) Use the `lm` function in R to fit the MLR model with *rating* as the response and the other 8 variables as predictors. Display the summary output.

(c)(3pts) Which predictor variables are statistically significant under the significance threshold value of 0.01?

(d)(2pts) What proportion of the total variation in the response is explained by the predictors?

(e)(3pts) What is the null hypothesis of the global F-test? What is the p-value for the global F-test? Do the 7 predictor variables explain a significant proportion of the variation in the response?

(f)(2pts) Consider testing the null hypothesis $H_0 : \beta_{carbo} = 0$, where β_{carbo} is the coefficient corresponding to *carbohydrates* in the MLR model. Use the t value available in the summary output to compute the p-value associated with this test, and verify that the p-value you get is identical to the p-value provided in the summary output.

(g)(4pts) Suppose we are interested in knowing if either *vitamins* or *potass* had any relation to the response *rating*. What would be the corresponding null hypothesis of this statistical test? Construct a F-test, report the corresponding p-value, and your conclusion.

(h)(3pts) Use the summary output to construct a 99% confidence interval for $\beta_{protein}$. What is the interpretation of this confidence interval?

(i)(3pts) What is the predicted *rating* for a cereal brand with the following information:

- Protein=3
- Fat=5
- Fiber=2
- Carbo=13
- Sugars=6
- Potass=60
- Vitamins=25
- Cups=0.8

(j). (3pts) What is the 95% prediction interval for the observation in part (i)? What is the interpretation of this prediction interval?

Q2.(20pts) Consider the MLR model with p predictors:

$$\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \boldsymbol{\epsilon}, \quad \boldsymbol{\epsilon} \sim N_n(\mathbf{0}, \sigma^2 \mathbf{I}_n)$$

If we define $\hat{\sigma}^2 = \frac{SSR}{n-p^*}$, with $p^* = p + 1$. Use theoretical results from the lectures to show that $\hat{\sigma}^2$ is an unbiased estimator of σ^2 . Find $V(\hat{\sigma}^2)$.